

Verifier-Gated and Auditable Multi-xApp Arbitration in the Near-RT RIC: Registered-Model Guarantees and the Limits of Incremental Repair

Lillia Ouali^{1*}, Madani Bezoui², Kamal Amroun¹, Ahcene Bounceur³

^{1*}Department of Computer Science, University of Bejaia, Bejaia, Algeria.

²LINEACT, CESI, Nancy, France.

³University of Sharjah, Sharjah, United Arab Emirates.

*Corresponding author(s). E-mail(s): lillia.ouali@univ-bejaia.dz;
Contributing authors: mbezoui@cesi.fr; kamal.amroun@univ-bejaia.dz;
abounceur@sharjah.ac.ae;

Abstract

Independently developed xApps may issue concurrent control requests that are mutually incompatible or violate registered service-level constraints. We formulate multi-xApp arbitration in the Near-RT RIC as a dynamic finite-domain constraint problem and present CONSISTXAPP, a verifier-gated pipeline that binds each decision to a registered model version and emits a replayable audit certificate (providing logical replayability without tamper-evidence unless cryptographic hashes are added) for every repair. Upstream, incremental generalized arc consistency removes unsupported values, while invalid candidates are repaired by solving expanding local subproblems guided by propagation explanations, solver infeasibility cores, or a deterministic graph frontier. Across four validation gates comprising 680,450 checks, the production verifier exhibited zero observed divergence from independently implemented references. On 250 generated static instances representing five O-RAN conflict scenarios, every verifier-accepted outcome satisfied all registered hard constraints. The latency benefit was amortized rather than absolute: under 2% sequential relation updates, repair of the previous accepted decision required median per-epoch latencies of 0.05–4.9 ms for 100–1,000 variables, compared with 2.8–25.6 ms for the evaluated Python CP-SAT rebuild-and-resolve baseline. The advantage narrowed near the pooled p95, reversed at the descriptive pooled p99, and disappeared on cold starts. Local repair also reduced objective quality relative to complete optimization. All guarantees are conditional on the correctness and completeness of the registered constraint model and do not establish physical-RAN safety or hard real-time schedulability.

Keywords: O-RAN, xApp conflict management, dynamic constraint satisfaction, generalized arc consistency, explanation-based repair, constraint programming, Near-RT RIC

1 Introduction

The Open Radio Access Network (O-RAN) architecture disaggregates the radio access network

and exposes its control through programmable components. Its near-real-time RAN Intelligent Controller (Near-RT RIC) hosts xApps, namely

third-party applications that subscribe to measurements and issue control requests through the E2 interface within control loops ranging from tens of milliseconds to one second [1]. Because xApps may be developed independently, often by different vendors, they may request incompatible values of the same parameter, steer distinct parameters toward jointly incompatible operating points, or collectively violate a service-level agreement (SLA). The O-RAN Alliance therefore assigns a conflict-mitigation function to the Near-RT RIC [2]. The associated literature commonly distinguishes direct, indirect, and implicit conflicts and proposes mechanisms for their detection and resolution [3, 4].

Most existing mechanisms primarily address what a candidate action may do to the network through profiling, digital twins, learning, or bargaining. This paper addresses the complementary *post-registration arbitration problem*: given finite action domains and registered hard relations, construct and audit a joint action such that no assignment violating the current registered model version is admitted to the E2 control path.

Throughout, *verified* means *accepted by the registered-model verifier*. It does not mean formally verified software: the paper provides architectural, differential, exhaustive, property-based, and mutation-based evidence for verifier correctness, not a machine-checked proof of the verifier implementation.

The following vocabulary is used consistently. A *proposal* is an individual xApp request. A *candidate* is a joint assignment before verification. An *accepted decision* is a candidate accepted by the verifier. A *forwarded action* is an accepted decision actually admitted to the E2 control path. A *certificate* is the replayable audit object associated with a repair.

We answer with CONSISTXAPP, a coordination pipeline that combines three established constraint-programming mechanisms in an O-RAN control architecture. First, an incremental generalized arc-consistency (GAC) stage removes values having no support in an incident hard relation, reuses propagation state across decision epochs, and restores affected values when a relation is relaxed. Second, when a decoded candidate violates one or more relations, a solver-based repair stage solves an induced subproblem. The violated constraints seed a variable core, variables outside the core are fixed to

the candidate values, and the core expands according to propagation explanations, solver infeasibility cores, or a deterministic frontier. The repair may escalate to the complete model. Third, an independently implemented verifier gates every forwarded action. The architecture is deliberately conservative: optimizer output is treated as a proposal until it has been checked against the current registered model.

The central contribution is a verifier-gated arbitration architecture with model-version binding and replayable audit certificates (providing logical replayability without tamper-evidence); incremental GAC and expanding local repair are upstream mechanisms for generating acceptable candidates efficiently. The paper does not claim a new consistency algorithm.

A complete modern constraint solver provides a strong baseline for the generated instances considered here. On the small static instances used in our scenario suite, rebuilding and solving the complete CP-SAT model requires less than 3 ms in the median. Consequently, the approximately 1 ms median difference in favour of CONSISTXAPP on this suite is not, by itself, a compelling reason to deploy an additional coordination layer. The intended operating regime is sequential: consecutive decision epochs are expected to share most of their variables, domains, and relations.

In the evaluated implementation, rebuilding a CP-SAT model at every epoch incurs model-construction and full-model search costs that scale with the complete generated instance, and supplying the previous solution as a hint did not materially reduce these costs. Repairing the previous accepted decision instead operated on the relations invalidated by the update and on the repair core induced by them. The scope of this comparison - an evaluated Python rebuild-and-resolve baseline, not every persistent or incremental solver architecture - is detailed in Section 8.7.

The contributions of this paper are as follows.

First, we implement and validate an independent forwarding verifier for registered multi-xApp arbitration. The verifier is evaluated through exhaustive comparison with an oracle, a second implementation, property-based testing, and oracle-labelled input mutation, with model-version binding tying every decision to the registered model

against which it was checked (Section 8 discusses the scope of this evidence).

Second, we make every non-trivial arbitration auditable: each repair emits a replayable certificate recording the violation witnesses, the core-expansion trace and its causes, and the final verifier outcome.

Third, we formulate registered arbitration as a dynamic finite-domain constraint problem with context-dependent hard relations, soft utility, and temporal-stability costs, and we integrate established constraint-programming mechanisms - incremental GAC with restoration, explanation-guided monotone core expansion with conditional completeness, and solver-based repair - into the verified pipeline. Expanding cores resolved all 248 evaluated repair candidates, whereas fixed violated-scope repair failed on 21 of them, this success-rate difference is statistically significant under the exact paired test. We also evaluate stronger conflict-aware decoders and show that they reduce, but do not eliminate, the need for verified repair.

Finally, we characterize the operating regime and its limitations: sequential repair provides a growing median advantage when updates are small, there is no cold-start advantage over complete CP-SAT up to 1,000 variables, the available sample supports only descriptive tail conclusions, local repair has a measurable objective-quality cost, and an evaluated continuous-relaxation stage is counterproductive.

Each contribution maps to a specific result, as summarized below.

- (C1) **Verifier-gated architecture:** Proposition 1 and validation gates G1–G4 (Section 4.5, Table 8).
- (C2) **Replayable audit certificate:** the certificate schema and validity conditions of Definition 1, with the worked replay of Section 3.3.
- (C3) **Expanding repair:** 248/248 candidates repaired by expanding cores against 227/248 under fixed violated-scope repair (Section 7.4).
- (C4) **Sequential operating regime:** median per-epoch latencies and their tail limitations (Section 7.8).
- (C5) **Negative boundaries:** absence of a cold-start advantage, the objective-quality gap,

and the counterproductive continuous relaxation (Sections 7.8 and 7.5).

These contributions answer six research questions stated in Section 6, each of which is addressed by a corresponding results subsection.

Section 2 reviews related work. Section 3 defines the coordination model. Section 4 presents CONSISTXAPP, and Section 5 gives its formal properties. Sections 6 and 7 describe the experimental setup and results. Section 8 discusses threats to validity, and Section 9 concludes the paper.

2 Background and Related Work

The Near-RT RIC supports control loops between tens of milliseconds and one second. xApps subscribe to key performance measurements (KPMs), process network context, and issue E2 control requests [1]. Experimental platforms such as CoLo-RAN, OpenRAN Gym, and FlexRIC support development and closed-loop evaluation of O-RAN controllers [5–7].

Three conflict classes are commonly distinguished [3]. A *direct conflict* occurs when two xApps request incompatible values of the same controlled parameter. An *indirect conflict* occurs when different controlled parameters influence a shared KPM. An *implicit conflict* couples distinct parameters and KPMs through network dynamics. Direct conflicts can often be checked syntactically. Indirect and implicit conflicts require a registered parameter–KPM dependency model, a validated profile, a simulator, or a conservative operator rule.

Existing proposals include a conflict-mitigation framework integrated into the RIC [3], QoS-aware optimization that maximizes the number of satisfied xApp requirements [4, 8], cooperative bargaining [9], pre-deployment profiling with statistical models and hierarchical graphs [10], digital-twin-based impact evaluation [11], and rule-based mitigation for mobility–energy interactions [12]. XAI4C applies explainable-AI techniques to conflict detection and mitigation [13]. Its explanations are post-hoc attributions of a learned mechanism, whereas the explanations considered here are logical premises replayable against a registered finite-domain model.

Table 1 positions CONSISTXAPP against the closest representative systems along the dimensions that matter for post-registration arbitration. The distinguishing combination is not any single column but the conjunction of a hard-feasibility gate, incremental repair, an audit trace, and model-version binding.

The comparison is deliberately narrow. These systems largely address what a candidate action *may do* to the network, a question CONSISTXAPP does not attempt to answer, conversely CONSISTXAPP enforces a registered model that those systems are, in part, mechanisms for constructing. The approaches are therefore complementary rather than competing, and the table should be read as a statement of scope, not of superiority.

On the constraint-programming side, propagation removes domain values that cannot participate in any locally supported tuple. For non-binary relations, this property is generalized arc consistency, which can be enforced efficiently for positive table constraints by simple tabular reduction and related algorithms [14, 15]. Local consistency is not a decision procedure: a GAC network may still be globally infeasible. It can nevertheless detect some contradictions and reduce the search space passed to a complete solver.

Dynamic constraint satisfaction studies sequences of related problems and the incremental maintenance of consistency. A line of work descending from DnAC-4 [16] maintains per-value justifications so that, after a constraint retraction, only a limited set of previously removed values must be reconsidered, AC|DC [17] and DnAC-6 [18] refine this idea with lower space and time overheads. The restoration rule used in this paper is deliberately coarser: after a relaxation, the affected constraint-graph component is reset and re-filtered (Section 4.1). This choice trades restoration minimality for simplicity of the audit trail, because component-level restoration is directly replayable by the explanation checker without per-value justification bookkeeping. We therefore position the propagation layer as systems integration rather than as an algorithmic advance over fine-grained dynamic arc-consistency maintenance, which we neither extend nor benchmark against. Explanation mechanisms record the causes of value removals and contradictions and support diagnosis and repair [19]. Dynamic backtracking similarly

preserves useful information while revising inconsistent decisions [20]. Large-neighbourhood search restricts a complete solver to a neighbourhood of a current assignment and modifies that neighbourhood as search proceeds [21]. Assumption-based infeasibility analysis identifies a sufficient subset of active assumptions causing inconsistency, such sets need not be cardinality-minimal.

Modern solvers such as OR-Tools CP-SAT combine presolve, propagation, clause learning, integer reasoning, and optimization [22, 23]. CONSISTXAPP does not claim to invent incremental consistency, explanations, assumption cores, or neighbourhood repair. Its contribution is their integration with an independently validated forwarding gate in registered xApp arbitration, together with an evaluation that identifies the sequential regime in which the integration was beneficial.

3 System Model

3.1 Epochs, actions, and registered relations

The Near-RT RIC operates over decision epochs $t = 1, \dots, T$. At epoch t , each active action variable $i \in \mathcal{A}_t$ has a finite base domain

$$D_i^t = \{a_{i,1}^t, \dots, a_{i,d_i}^t\}. \quad (1)$$

A value may encode acceptance or rejection of a request, a target cell, an execution slot, a discrete power level, a handover offset, a resource-allocation profile, or a registered no-operation. The coordinated action is $x_i^t \in D_i^t$. The time index is omitted when a single epoch is considered.

A constraint $c \in \mathcal{C}_t$ has a scope $S_c \subseteq \mathcal{A}_t$ and an allowed relation

$$R_c(\mathbf{s}_t) \subseteq \prod_{i \in S_c} D_i^t, \quad (2)$$

where $\mathbf{s}_t$ denotes the registered context. An assignment $\mathbf{x}^t$ satisfies c when

$$\mathbf{x}_{S_c}^t \in R_c(\mathbf{s}_t). \quad (3)$$

Relations may be binary or higher-order and may be derived from operator policies, ownership rules, SLA thresholds, validated xApp profiles, simulators, digital twins, or conservative safety

Table 1 Positioning of CONSISTXAPP against representative conflict-management systems. Entries reflect the mechanisms emphasized in the cited descriptions, “Not reported (n/r)” denotes a property not reported as a design element of that work rather than a demonstrated absence.

Work	Conflict classes	Model basis	Hard-feas. gate	Incr. repair	Audit trace	Version binding
QACM [4]	Direct and indirect	QoS-aware request optimization	Not reported (n/r)	Not reported (n/r)	Not reported (n/r)	Not reported (n/r)
Bargaining [9]	Direct and indirect	Cooperative game-theoretic	Not reported (n/r)	Not reported (n/r)	Not reported (n/r)	Not reported (n/r)
PACIFISTA [10]	Direct, indirect, implicit	Pre-deployment statistical profiling	Not reported (n/r)	Not reported (n/r)	Profiles	Not reported (n/r)
COMIX [11]	Indirect and implicit	Digital-twin impact evaluation	Not reported (n/r)	Not reported (n/r)	Not reported (n/r)	Not reported (n/r)
XAI4C [13]	Direct and indirect	Learned model with post-hoc XAI	Not reported (n/r)	Not reported (n/r)	Attributions	Not reported (n/r)
CONSISTXAPP	Direct, indirect, implicit when registered	Registered finite-domain model	Yes	Yes	Yes	Yes

rules. Hard constraints must be satisfied before forwarding. Soft constraints contribute only to utility.

Let $\mathcal{C}_{t,\text{hard}} \subseteq \mathcal{C}_t$ denote the hard constraints. The registered hard model is

$$\mathcal{M}_t = (\mathcal{A}_t, \mathbf{D}_t, \mathcal{C}_{t,\text{hard}}, \mathbf{s}_t, \nu_t), \quad (4)$$

where ν_t is a model-version identifier. A verification result is valid only for the model version against which it was obtained. A production implementation must therefore bind the decision, certificate, and forwarding event to the same version and re-verify a candidate if the registered model changes before forwarding.

Between epochs, xApps may activate or deactivate, domains and constraints may change, relations may tighten or relax, and SLA thresholds may move. The modified elements form a change set Δ_t . Restrictive and relaxing changes must be distinguished. After a restriction, previous support information may remain useful. After a relaxation, a previously removed value may have regained support and must be reconsidered.

3.2 Utility and temporal stability

Let $u_i(a_i, \mathbf{s}_t)$ be the utility of choosing action a_i for variable i , with weight $\omega_i \geq 0$. Let $v_q \geq 0$ be a soft violation with weight $\lambda_q \geq 0$. The static utility is

$$U(\mathbf{x}, \mathbf{s}_t) = \sum_{i \in \mathcal{A}_t} \omega_i u_i(x_i, \mathbf{s}_t) - \sum_{q \in \mathcal{Q}_t} \lambda_q v_q(\mathbf{x}, \mathbf{s}_t). \quad (5)$$

Given the previous accepted decision $\mathbf{x}^{t-1}$ and reconfiguration costs $\kappa_i \geq 0$, the weighted action

churn is

$$H(\mathbf{x}^t, \mathbf{x}^{t-1}) = \sum_{i \in \mathcal{A}_t \cap \mathcal{A}_{t-1}} \kappa_i \mathbf{1}[x_i^t \neq x_i^{t-1}]. \quad (6)$$

The general scalarized coordination problem is

$$\begin{aligned} \max_{\mathbf{x}^t} \quad & U(\mathbf{x}^t, \mathbf{s}_t) - \mu H(\mathbf{x}^t, \mathbf{x}^{t-1}) \\ \text{s.t.} \quad & x_i^t \in D_i^t, & i \in \mathcal{A}_t, \\ & \mathbf{x}_{S_c}^t \in R_c(\mathbf{s}_t), & c \in \mathcal{C}_{t,\text{hard}}, \end{aligned} \quad (7)$$

where $\mu \geq 0$ controls the preference for temporal stability. In this formulation, utility can compensate for additional churn depending on μ . However, the operational pipeline proposed in this paper implements a strictly lexicographic repair policy (Section 4.3) where no utility improvement can compensate for any increase in primary churn. The complete-objective baseline evaluated in Section 7.5 optimizes the lexicographic formulation using the complete CP-SAT model directly on the raw state. Dwell-time and cooldown rules are represented as hard relations when their violation is not permitted.

3.3 Worked example

Consider a coverage xApp controlling

$$x_1 \in \{\text{low, medium, high}\}$$

and an energy-saving xApp controlling

$$x_2 \in \{\text{sleep, on}\}.$$

Table 2 Principal notation used throughout the paper.

Symbol	Meaning
t, T	Decision epoch index and horizon
$\mathcal{A}_t$	Active action variables at epoch t
D_i^t, d_i	Finite base domain of variable i and its size
$x_i^t, \mathbf{x}^t$	Coordinated action value and joint assignment
$\mathcal{C}_t, \mathcal{C}_{t,\text{hard}}$	Constraints and hard constraints
$c, S_c, R_c(\mathbf{s}_t)$	Constraint, its scope, its allowed relation
$\mathbf{s}_t$	Registered context
$\mathcal{M}_t, \nu_t$	Registered hard model and its version identifier
Δ_t	Change set between epochs $t-1$ and t
u_i, ω_i	Per-variable utility and its weight
$v_q, \lambda_q, \mathcal{Q}_t$	Soft violation, weight, and soft-constraint set
H, κ_i, μ	Action churn, reconfiguration cost, stability weight
U	Static utility of an assignment
$\mathcal{A}_r$	Repair core (variables reoptimized during repair)
$\mathcal{X}_t$	Verifier input representation space
$\mathcal{F}(\mathbf{s}_t)$	Feasible set of the registered hard model

Suppose the registered relation is

$$R_{12} = \{(\text{medium}, \text{on}), (\text{high}, \text{on}), (\text{low}, \text{sleep})\}. \quad (9)$$

The relation forbids high power when the cell is sleeping and low power when the cell is active. Assume that the preferred requests are $x_1 = \text{high}$ and $x_2 = \text{sleep}$. These requests are soft preferences, the base domains remain multivalued.

GAC removes no value. The value high has support (high, on), and sleep has support (low, sleep). Nevertheless, the independently decoded pair (high, sleep) violates R_{12} . The verifier rejects it, the violated relation seeds the repair core $\{x_1, x_2\}$, and repair may return (high, on) or (low, sleep) depending on the registered preferences. Thus propagation removes individually unsupported values, whereas repair resolves incompatible combinations of individually supported values.

Certificate emitted by this repair

Assuming the registered preferences select (high, on), the arbitration above emits the following certificate:

```
model_version:    nu17
candidate:        (high, sleep)
```

```
violated_relation: R12
initial_core:      {x1, x2}
expansion_trace:   none
final_assignment:  (high, on)
solver_status:     OPTIMAL
verifier_status:   ACCEPT
```

Replaying this certificate under Definition 1 requires checking that the recorded initial violation is reproducible and that the final assignment is feasible, that is,

$$(\text{high}, \text{sleep}) \notin R_{12}, \quad (\text{high}, \text{on}) \in R_{12}, \quad (10)$$

both of which hold by inspection of R_{12} . The empty expansion trace is consistent because the initial core already spanned every variable of the violated relation, so no expansion was required. An auditor can perform these checks without access to the solver, the decoder, or the propagation engine.

4 The ConsistXApp Framework

Figure 1 positions CONSISTXAPP in the Near-RT RIC control path. The pipeline separates four stages whose boundaries are maintained throughout the paper:

1. **Candidate generation:** incremental GAC followed by a decoder, or, in the sequential regime, the previous accepted decision.
2. **Repair feasibility:** expansion of a repair core solved with objective-free CP-SAT, establishing hard feasibility only.
3. **Operational optimization:** sequential lexicographic optimization restricted to the retained core.
4. **Independent verification:** a full registered-model membership test on the resulting candidate, followed by the forwarding gate.

Keeping these stages distinct avoids conflating four different claims: finding a feasible core, optimizing within that core, certifying optimality, and verifying hard feasibility. Only the fourth stage supports the forwarding guarantee.




Fig. 1 CONSISTXAPP in the Near-RT RIC control path. Solid arrows carry control proposals and accepted actions, dashed arrows carry registered-model and context data, dotted arrows carry audit data. Every forwarded action passes the independent verifier V_t and an atomic model-version check bound to ν_t : a candidate rejected by the verifier, or one whose model version changed between verification and forwarding, is not admitted to E2. Incremental GAC and solver-based repair generate candidates efficiently but are outside the trusted path - a fault in either can reduce availability or objective quality, but cannot by itself cause a registered-model-invalid forwarding.

4.1 Incremental consistency filtering

The propagator maintains a base domain D_i and a filtered domain $D_i^* \subseteq D_i$ for every variable. Keeping them separate is necessary because a later relaxation may restore support for a value that was previously removed.

For $a \in D_i$ and constraint c , the support set is

$$\text{Supp}_c(i, a; \mathbf{D}) = \{ \mathbf{z} \in R_c(\mathbf{s}_t) : z_i = a \wedge z_j \in D_j, \forall j \in S_c \}. \quad (11)$$

A hard network is generalized arc consistent when every value of every variable has at least one support in every incident relation. The revision operator is

$$\begin{aligned} \text{Revise}_{c,i}(\mathbf{D}) : \\ D_i \leftarrow \{ a \in D_i : \text{Supp}_c(i, a; \mathbf{D}) \neq \emptyset \}. \end{aligned} \quad (12)$$

Revision is applied through a work queue until the fixed point

$$\mathbf{D}^* = \text{GAC}(\mathbf{D}, \mathcal{C}_{\text{hard}}, \mathbf{s}_t) \quad (13)$$

is reached. An empty filtered domain is a *wipeout* and proves that no feasible assignment extends the current domains. Non-empty filtered domains do not prove global feasibility.

Across epochs, the propagator is incremental. The queue is seeded with the constraints affected by Δ_t and constraints incident to an affected variable. Under purely restrictive changes, cached supports and previous deletions can be reused when their premises remain valid. If the update contains a

relaxation, the affected constraint-graph component is restored to its base domains and current allowed tuples before GAC is re-established. This full-component restoration is a conservative sufficient rule, it is not claimed to be the minimum restoration set computed by fine-grained dynamic arc-consistency schemes [16–18], and is preferred here because the restored set is replayable by the explanation checker without per-value justification state.

For explicit positive tables, the implementation records active tuples, the last known support of each variable–value pair, the context identifier under which the support was valid, and the explanation attached to each removal.

4.2 Propagation explanations

When a value a is removed from D_i by constraint c , the propagator records an explanation

$$\mathcal{E}(i, a) \subseteq \mathcal{C}_{\text{hard}} \cup \mathcal{L}, \quad (14)$$

where $\mathcal{L}$ contains temporary domain restrictions and boundary decisions.

The explanation contains premises sufficient to establish that every tuple of R_c assigning a to x_i is inactive. For each candidate support tuple, at least one recorded premise explains why the tuple cannot be used. A wipeout explanation is

$$\mathcal{E}_\emptyset(i) = \bigcup_{a \in D_i} \mathcal{E}(i, a). \quad (15)$$

Algorithm 1 Incremental GAC with restoration after relaxation

Input: warm filtered domains D_{t-1}^* , base domains, active tables, and change set Δ_t

Output: filtered domains D_t^* or WIPEOUT

- 1: Classify Δ_t as restrictive, relaxing, mixed, or utility-only.
 - 2: If Δ_t contains a relaxation, restore the affected constraint-graph component to its base domains and current allowed tuples.
 - 3: Apply valid restrictive deletions to the warm domains.
 - 4: Seed the work queue with changed constraints and constraints incident to changed variables.
 - 5: While the queue is non-empty, remove a constraint c and revise each $i \in S_c$.
 - 6: If some revised domain becomes empty, return WIPEOUT with its explanation.
 - 7: Whenever a domain shrinks, enqueue the constraints incident to that domain.
 - 8: Return the fixed-point domains D_t^* .
-

Explanations are not required to be cardinality-minimal. They must be valid relative to the registered model, reproducible from the propagation trace, associated with a model/context identifier, and accepted by a separate explanation checker. The checker replays the premises and verifies that every potential supporting tuple is disabled.

4.3 Decoding and explanation-guided solver-based repair

If propagation terminates without a wipeout, a deterministic decoder selects

$$\hat{x}_i \in D_i^* \quad (16)$$

as the highest-utility value under a documented tie-breaking rule. This utility decoder is intentionally conflict-unaware and is used to isolate the contribution of repair. Section 7.7 separately evaluates conflict-aware candidate generators.

Every decoded candidate is checked against all hard relations. Let

$$\mathcal{C}^0 = \{c : \hat{x}_{S_c} \notin R_c(s_t)\} \quad (17)$$

be the initially violated relations. The initial variable core is

$$\mathcal{A}^0 = \bigcup_{c \in \mathcal{C}^0} S_c. \quad (18)$$

At repair iteration r , variables outside $\mathcal{A}^r$ are fixed to their decoded values, while variables inside the core retain their filtered domains. GAC is enforced on the restricted model before the sub-problem is submitted to CP-SAT. A variable $j \notin \mathcal{A}^r$ that participates in a constraint whose scope intersects $\mathcal{A}^r$ is called a *boundary variable* of iteration r .

Two further objects are used by the expansion rule. Let $\text{VarLit}(\mathcal{E})$ denote the set of variables appearing in the temporary restriction literals or boundary decisions contained in an explanation $\mathcal{E}$, and let $\mathcal{E}_\emptyset^r$ denote the wipeout explanation recorded at iteration r , that is, the set of literals whose conjunction entails the empty domain. The set $\mathcal{C}^{r+1}$ denotes the relations found violated or wiped out at iteration r and used to seed the next core.

The *deterministic frontier rule* is defined as follows. Constraints are indexed once, at model registration, by a fixed total order (ascending registered constraint identifier). Adjacency is taken through the constraint hypergraph, in which two variables are adjacent when they co-occur in the scope of some active hard relation. At each frontier expansion, the rule adds exactly one constraint: the lowest-indexed active hard relation whose scope intersects $\mathcal{A}^r$ but is not contained in it, together with all variables in that scope. Ties cannot arise because the constraint index is a total order. One constraint, not one adjacency layer, is added per iteration, this keeps expansions minimal and makes certificate replay deterministic.

If propagation produces a wipeout, its explanation supplies constraints and variables for expansion:

$$\mathcal{A}^{r+1} = \mathcal{A}^r \cup \text{VarLit}(\mathcal{E}_\emptyset^r) \cup \bigcup_{c \in \mathcal{C}^{r+1}} S_c, \quad (19)$$

where

$$\mathcal{C}^{r+1} = \mathcal{E}_\emptyset^r \cap \mathcal{C}_{\text{hard}}. \quad (20)$$

If propagation succeeds but the restricted model is reported **INFEASIBLE**, the assumption-core variant uses reified boundary fixings. For each

boundary variable j , a Boolean assumption literal b_j enforces

$$b_j \Rightarrow (x_j = \hat{x}_j). \quad (21)$$

The literals returned by `SufficientAssumptionsForInfeasibility` form a sufficient, not necessarily minimum, set of assumptions whose simultaneous activation is inconsistent with the remaining model. The corresponding boundary variables are released. If the returned set is empty, all currently fixed boundary variables are conservatively released. A solver status of UNKNOWN is not interpreted as infeasibility.

The core-extraction model is a feasibility model and contains no objective. In the default explanation-guided variant, boundary values are represented by equality constraints. If neither an explanation nor an assumption core adds a variable, a deterministic frontier rule adds adjacent constraints and their variables.

Algorithm 2 Explanation-guided repair-core expansion

Input: decoded candidate $\hat{x}$, registered hard model, stage-boundary experimental budget, and the registered fallback associated with the current model version

Output: verified local repair, verified full-scope repair, verified registered fallback, or BLOCKED

- 1: **Core feasibility phase:** Compute the violated relation set $\mathcal{C}^0$ and set $\mathcal{A}^0 \leftarrow \bigcup_{c \in \mathcal{C}^0} S_c$.
 - 2: For repair iteration $r = 0, 1, \dots$, fix each boundary variable $j \notin \mathcal{A}^r$ to $\hat{x}_j$.
 - 3: Enforce GAC on the restricted model.
 - 4: If propagation wipes out at the full core $\mathcal{A}^r = \mathcal{A}_t$, the registered hard model is infeasible, no same-model candidate, including the registered fallback, can pass the verifier, so return BLOCKED for the current model version.
 - 5: If propagation wipes out on a proper core $\mathcal{A}^r \subsetneq \mathcal{A}_t$, extract the wipeout explanation, expand the core, and continue.
 - 6: Otherwise solve the objective-free restricted subproblem with CP-SAT.
 - 7: If the solver reports INFEASIBLE at the full core $\mathcal{A}^r = \mathcal{A}_t$, the registered hard model is infeasible under the current exact encoding and conclusive status, return BLOCKED for the current model version.
 - 8: If the solver reports INFEASIBLE on a proper core, expand the core using an assumption core when available, otherwise use the deterministic frontier rule, and continue.
 - 9: If the solver returns UNKNOWN, no infeasibility conclusion is drawn, invoke the registered fallback policy (since an UNKNOWN status typically indicates that the solver budget or deadline has expired) and go to the verification phase with the fallback candidate.
 - 10: If the solver returns FEASIBLE or OPTIMAL, retain the current core $\mathcal{A}^r$ and break to the optimization phase.
 - 11: **Operational optimization phase:** Optimize the lexicographic objectives on the feasible core. If every active lexicographic level returns OPTIMAL, mark the repair LEX-OPTIMAL. If the remaining optimization budget expires after the solver has produced a feasible incumbent but before all lexicographic levels are certified, mark the result FEASIBLE-NOT-CERTIFIED and submit the incumbent to the independent verification phase. The incumbent is not treated as verified at this point, independent verification occurs only in Step 12.
 - 12: **Verification phase:** Submit the final optimized or fallback candidate to the independent verifier. If the verifier accepts it, return the verified candidate. If it rejects an optimized candidate, submit the registered fallback to the verifier, if the fallback is accepted, return it. Otherwise return BLOCKED.
-

Each non-trivial arbitration emits a *replayable audit certificate* containing sufficient information to reconstruct the initial violations, validate each recorded core expansion, and recheck final registered-model feasibility. Concretely, the certificate records the model and context identifiers, the decoded candidate, the initial violation witnesses, the repair-core sequence, the source of every expansion, the released boundary variables, the final assignment, the solver status, and the verifier outcome. Propagation explanations and solver assumption cores are included when they cause an expansion. The certificate does not prove minimum-core size, solver correctness, or global objective optimality.

Definition 1 (Certificate validity). *A certificate π is valid for the registered model $\mathcal{M}_t$ at version ν_t when all of the following hold:*

- (1) *the model and context identifiers recorded in π match $(\mathcal{M}_t, \nu_t)$;*
- (2) *the listed initial violations are reproducible by evaluating the recorded candidate against the registered relations;*
- (3) *every recorded expansion follows the declared explanation, assumption-core, or deterministic frontier rule;*
- (4) *all released boundary variables are recorded;*
- (5) *the final assignment is complete and domain-valid;*
- (6) *every registered hard relation is satisfied by the final assignment;*
- (7) *the recorded verifier outcome and the forwarded assignment agree.*

Certificate validity establishes replayable registered-model feasibility and trace consistency. It does not establish solver optimality, physical-network safety, or completeness of the registered model. The present evaluation did not separately measure serialized certificate size or replay latency.

The intended repair preference order is lexicographic:

1. satisfy every hard relation;
2. minimize weighted reconfiguration against the previous verified decision;
3. minimize deviation from the decoded candidate;
4. maximize static utility.

Table 3 Permitted conclusions per solver status.

Status	Permitted conclusion
OPTIMAL	Feasible and optimal for the current encoded objective
FEASIBLE	Feasible incumbent, optimality not established
INFEASIBLE	Infeasible for the current encoded model and boundary conditions only
UNKNOWN	No feasibility or infeasibility conclusion

Hard feasibility is not scalarized with the soft levels. A **FEASIBLE** solver status establishes feasibility but does not establish objective optimality. A claim of global objective optimality requires a corresponding **OPTIMAL** status for the encoded complete model. Accordingly, this paper uses the term *solver-based repair* for the CP-SAT-based repair engine. The evaluation therefore reserves *certified optimum* (qualified by the exact encoded model, integer scaling of utilities, and active lexicographic levels) for solutions proven optimal across all active lexicographic levels. Table 3 states the conclusion permitted by each solver status.

In particular, **INFEASIBLE** on a proper repair core does not imply that the complete registered model is infeasible: the fixed boundary may exclude assignments that the complete model admits.

Fallback semantics.

A fallback is a version-specific candidate policy, not an exemption from verification. Registration is mandatory: a fallback must be explicitly registered in the current model version - for example, rejection of the conflicting requests, retention of a previous configuration, or a safe no-operation - and fallback lookup is performed against the active version ν_t . It must pass the same verifier as any solver-generated candidate, so a previous configuration is re-verified before reuse. Fallback validation is not cached across model versions, any change of ν_t invalidates a prior acceptance, because the relations against which it was checked may have changed. If no fallback is registered, or if the registered fallback fails verification under the current model version, the outcome is **BLOCKED**. If the complete registered hard model is infeasible, no fallback belonging to that model can be accepted,

and the pipeline returns BLOCKED for that model version.

4.4 Independent verification

The production verifier is implemented separately from propagation, explanation construction, and repair. It receives the model/context identifier, selected actions, base domains, hard relations, and SLA thresholds. It accepts the registered feasible set

$$\mathcal{F}(s_t) = \{x : x_i \in D_i \wedge x_{S_c} \in R_c(s_t), \forall c \in C_{\text{hard}}\}. \quad (22)$$

Proposition 1 (Registered-model forwarding soundness). *Let*

$$V_t : \mathcal{X}_t \longrightarrow \{\text{ACCEPT}, \text{REJECT}\} \quad (23)$$

be the production verifier bound to model version ν_t , where $\mathcal{X}_t$ is the space of candidate input representations submitted to the gate. The well-formed complete assignments $\prod_{i \in \mathcal{A}_t} D_i^t$ form a subset of $\mathcal{X}_t$; the remaining elements of $\mathcal{X}_t$ are malformed or incomplete submissions, namely those with missing variables, out-of-domain values, or stale context identifiers, all of which V_t maps to REJECT. If (i) $V_t(x) = \text{ACCEPT} \iff x \in \mathcal{F}(s_t)$ for every $x \in \mathcal{X}_t$, that is, V_t implements set membership in $\mathcal{F}(s_t)$ over its full input space; (ii) forwarding is enabled only after verifier acceptance, and (iii) the registered model version does not change between verification and forwarding, then every forwarded action belongs to $\mathcal{F}(s_t)$ for model version ν_t .

Proof The verifier checks the membership of every selected value in its base domain and checks every registered hard relation, so acceptance under (i) is equivalent to membership in $\mathcal{F}(s_t)$. The forwarding gate (ii) rejects any candidate that fails one of these checks, and model-version binding (iii) prevents acceptance under one registered model from being reused after the model changes. $\square$

Remark 1. *Premise (i) is the only premise about program behaviour rather than architecture, and it is precisely the property targeted empirically by validation gates G1–G4 (Section 4.5): G1 checks V_t against exhaustive enumeration of $\mathcal{F}(s_t)$ on small instances, G2 checks agreement with two*

independent implementations of the same membership predicate, and G3–G4 probe invariances and mutation behaviour that any membership test must satisfy. Proposition 1 is therefore an architectural invariant conditional on verifier correctness, complete registration of the hard constraints, and atomic model-version handling, the gates provide evidence for, not proof of, premise (i), and nothing here implies that the registered model captures every dependency of the physical RAN.

Model-version transaction.

Premise (iii) of Proposition 1 is an assumption about atomicity, and a production implementation must discharge it operationally rather than assume it. Model lookup, verifier execution, acceptance recording, and forwarding authorization must be treated as a single version-bound transaction. If the active version changes before forwarding, the authorization token is invalidated and the candidate is re-verified. The required sequence is:

- (1) load the active model version ν_t ;
- (2) generate a candidate under ν_t ;
- (3) verify the candidate under ν_t ;
- (4) atomically compare the current active version with ν_t ;
- (5) forward if the comparison succeeds, otherwise discard the authorization and restart from step (1).

This makes the time-of-check-to-time-of-use requirement operational. Without step (4), a model update committed between verification and forwarding could admit an action that the updated model rejects.

4.5 Verifier-validation gates

The verifier is evaluated through four validation gates.

- **G1: exhaustive oracle comparison.** Every assignment of randomly generated small instances is enumerated and classified by an independently coded oracle.
- **G2: triple agreement.** A declarative CP-SAT model is reconstructed from a JSON-normalized representation using a separate parser. The production verifier, oracle, and declarative implementation must agree.

- **G3: property-based testing.** The tests cover determinism, constraint-order invariance, variable-renaming invariance, and rejection of out-of-domain values, missing variables, and stale contexts.
- **G4: oracle-labelled input mutation.** Mutants of planted-feasible assignments are classified independently by the oracle. A mutant that remains feasible is labelled feasible; the test does not assume that every mutation must create an invalid assignment.

The oracle and second verifier share no relation parser, tuple normalizer, domain-indexing routine, or context-validation routine with the production verifier. A separate serializer produces the normalized representation consumed by G2. A divergence at any gate blocks the corresponding campaign. Zero divergence increases confidence in the tested implementation but does not prove correctness for all possible inputs.

4.6 Safety, availability, and timeliness

Three properties must be distinguished.

Safety means that no candidate rejected by the registered-model verifier is forwarded. *Availability* means that the pipeline produces a verifier-accepted action, including an accepted fallback. *Timeliness* means that such an action becomes available before the control deadline.

The forwarding architecture enforces safety relative to the registered model. Availability is not unconditional because the registered hard model or every fallback may be infeasible. The temporal metrics distinguish four boundaries: (1) CP-SAT’s internal 100 ms solver time limit, (2) stage-boundary budget checks, where an algorithm inspects elapsed time before invoking a subsequent stage, (3) post-hoc deadline-compliance classification, which labels an outcome based on total elapsed time, and (4) the absence of pre-emptive interruption, meaning a running stage (such as propagation) is allowed to complete before its latency is recorded. Consequently, the experiments characterize computational latency and post-hoc compliance, they do not demonstrate a hard real-time scheduling guarantee.

5 Analysis

The following properties of the propagation and repair layers are standard consequences of GAC theory and of the monotone expansion rule, we state them together with proof sketches because the pipeline relies on them, without claiming analytical novelty.

Proposition 2 (GAC solution preservation). *For every $\mathbf{x} \in \mathcal{F}(\mathbf{s}_t)$ and every $i \in \mathcal{A}_t$,*

$$\mathbf{x} \in \mathcal{F}(\mathbf{s}_t) \implies x_i \in D_i^*, \quad \forall i \in \mathcal{A}_t. \quad (24)$$

That is, GAC filtering removes no value used by a globally feasible assignment.

Proof sketch A value is removed only if it has no support under the current domains. By induction over the revision sequence, no component of a feasible assignment $\mathbf{x}$ can be the first removed component, because the remaining components of the same feasible tuple would still support it. $\square$

Proposition 3 (Monotone repair-core expansion). *The repair cores generated by Algorithm 2 satisfy*

$$\mathcal{A}^0 \subseteq \mathcal{A}^1 \subseteq \dots \subseteq \mathcal{A}_t. \quad (25)$$

If every unsuccessful non-full iteration adds at least one variable, the procedure performs at most $|\mathcal{A}_t| - |\mathcal{A}^0|$ strict expansions before reaching the full core.

Proof sketch Monotonicity and the expansion bound follow from the expansion rule, which only adds variables and never removes them, each strict expansion increases the core cardinality by at least one, and the core is bounded above by $\mathcal{A}_t$. $\square$

Proposition 4 (Conditional completeness at full scope). *Assume that (i) all domains are finite; (ii) every registered hard relation is encoded exactly; (iii) no active hard constraint is omitted from the encoding; (iv) the full core $\mathcal{A}^r = \mathcal{A}_t$ is reached; (v) the solver is complete for the encoded model, and (vi) the solver returns a conclusive status before its time limit. Then the repair stage returns a feasible assignment if and only if the registered hard model is feasible.*

Proof sketch At the full core no boundary variable is fixed, so the repair model coincides with the complete registered hard model. The claim then follows from assumptions (v) and (vi). Each assumption is necessary: relaxing (iii) or (vi) breaks the “only if” direction, and relaxing (iv) leaves a fixed boundary that may render a globally feasible model locally infeasible. $\square$

Proposition 5 (Necessity of restoration after relaxation). *If value a was removed at epoch $t - 1$ for lack of support in $R_c(s_{t-1})$ and the relation is relaxed at epoch t , then retaining the deletion without re-evaluation may exclude a feasible assignment of $\mathcal{M}_t$.*

Proof sketch The relaxed relation may introduce a tuple $z \in R_c(s_t)$ with $z_i = a$ whose other components belong to their current domains. Then z is a new support for a , so the previous no-support explanation is no longer valid and the deletion is unjustified. $\square$

Filtering alone does not guarantee valid decoding. Consider

$$D_1 = D_2 = \{0, 1\}$$

with

$$R_{\neq} = \{(0, 1), (1, 0)\}.$$

Every value has support, so the network is GAC. A deterministic per-variable decoder that resolves both symmetric ties identically returns either $(0, 0)$ or $(1, 1)$, both of which violate $R_{\neq}$. Propagation and repair therefore address different failure modes.

A repair restricted to a proper subset is not complete. Its infeasibility may result from the fixed boundary even if the complete model is feasible. Therefore, local infeasibility is never reported as global model infeasibility.

A complete scan of explicit positive tables admits the coarse per-scan bound

$$O\left(\sum_c |R_c| |S_c|\right). \quad (26)$$

Actual propagation time depends on support caching, tuple activity, arity, and the number of removals. Memory use includes the positive tables, active-tuple state, support caches, explanations, and certificates. Solver-based repair remains exponential in the worst case. Core expansion may

reach the complete variable set, at which point the localization advantage disappears.

Research questions

The evaluation is organized around six research questions, each answered by a designated results subsection.

RQ1 *Forwarding correctness evidence.* Does the independently implemented verifier agree with independent references across exhaustive, differential, property-based, and mutation-based tests? (Section 7.1)

RQ2 *Repair effectiveness.* Does expanding the repair core recover feasible decisions that fixed violated-scope repair misses? (Section 7.4)

RQ3 *Localization benefit.* How do explanation-guided cores compare with frontier, assumption-core, and full-scope repair in success rate, core size, and latency? (Section 7.4)

RQ4 *Operating regime.* Under what update rates and problem sizes does incremental repair outperform rebuild-and-resolve CP-SAT? (Section 7.8)

RQ5 *Cost of localization.* What objective-quality loss results from restricting optimization to a local repair core? (Section 7.5)

RQ6 *Boundary conditions.* How do cold starts, dense instances, large repair cores, and tail events affect the observed advantage? (Section 7.8)

6 Experimental Setup

6.1 Simulator, scenarios, and instance suites

The evaluation uses a Python-based O-RAN coordination simulator containing instance generators, an observable GAC propagator, explanation recorder and checker, OR-Tools CP-SAT adapters, the repair engine, the independent verifier, and a pseudo-physical network model.

The pseudo-physical model maps power, resource-block fraction, handover offset, and sleep-state decisions to SINR, throughput, energy per bit, Jain fairness, and a handover-failure indicator. Users are placed uniformly and path loss is

distance based. The model is intended to provide reproducible KPMs, not to replace a system-level or closed-loop O-RAN simulator.

Generated relations encode KPM thresholds, ownership rules, and sleep/wake exclusivity. Each generated feasible test instance retains a planted feasible assignment so that coordination failure can be distinguished from instance infeasibility. A separate diagnostic suite contains controlled infeasible instances for wipeout and explanation tests.

Table 4 summarizes the five scenarios. They contain 12–30 action variables with domain sizes of three or four and 50 seeds per scenario, giving 250 matched static instances.

The dynamic suite runs the five scenarios over 30 episodes of 100 epochs, giving 15,000 epochs. A fixed schedule contains restrictive, relaxing, mixed, and utility-only transitions that modify 1–50% of the model. Mixed transitions account for 84% of the schedule. Restrictive, relaxing, and utility-only transitions are included as isolated diagnostic cases.

The repair-stress suite contains 50 planted-feasible instances with 24 variables, domain size four, 34 relations, and 70% equality-type constraints. The candidate is generated by flipping three variables of the planted solution.

The cold-start suite contains instances with

$$n \in \{30, 100, 300, 600, 1000\}, \quad d = 6, \quad |\mathcal{C}| = 2n,$$

table tightness 0.25, and 20 seeds per size.

The sequential suite contains ten episodes of 30 epochs per size and change rate. Between epochs, a fraction

$$\delta \in \{2\%, 10\%\}$$

of the relations is re-randomized while retaining a planted feasible assignment. A drift variant replants 1% of the variables per epoch.

These generators produce controlled finite-domain workloads. They do not establish representativeness for dense, high-arity, adversarial, or commercial multi-vendor RAN models.

6.2 Operational repair objective and solver configuration

The core-extraction model focuses exclusively on feasibility to isolate conflicts, containing no objective function. The subsequent operational repair stage implements the lexicographic preference defined in Section 4.3 via sequential solves. We determine the optimum for each objective level sequentially:

$$z_1^* = \min_{x \in \mathcal{F}(s_t)} H(x, x^{t-1}) \quad (27)$$

$$z_2^* = \min_{\substack{x \in \mathcal{F}(s_t) \\ H(x, x^{t-1}) = z_1^*}} D(x, \hat{x}) \quad (28)$$

$$z_3^* = \max_{\substack{x \in \mathcal{F}(s_t) \\ H = z_1^* \\ D = z_2^*}} U_{\text{int}}(x, s_t) \quad (29)$$

where H is the weighted reconfiguration distance, D is the deviation from the decoded candidate $\hat{x}$, and U_{int} is the scaled static utility.

To comply with CP-SAT’s exact-integer requirements, continuous utility values and fractional weights are scaled by 10^4 and rounded to integers. The total 100 ms solver time allocation is shared sequentially: each CP-SAT call receives the remaining solver budget $B_{k+1} = \max\{0, B_k - \tau_k\}$, rather than the original full budget. The implementation applies the objective levels sequentially. A subsequent level is optimized only after the preceding level has returned **OPTIMAL**, the certified optimum value is then imposed as an equality constraint. If a stage returns **FEASIBLE** because its budget expires, the incumbent may be accepted after independent verification, but neither that objective level nor the complete lexicographic sequence is considered optimal. Such outcomes are reported as feasible repairs without an objective-optimality certificate. For tie-breaking, the utility decoder selects the highest utility, then lowest action index. For solver-generated repairs, tied optima are considered equivalent and any tied solution may be returned.

Table 5 lists the execution environments. All campaigns enforce single-threaded solver execution (`num_search_workers = 1`) with deterministic instance seeds, eliminating portfolio-search variance and making instance generation reproducible. The experiment harness itself is sequential: no campaign uses process-level parallelism, so the

Table 4 Generated O-RAN coordination scenarios.

Scenario	xApps	Controlled parameters	Principal KPMs	Conflict mechanism
S1: direct power conflict	Coverage and energy-saving	Transmission power, cell state	Coverage, SINR, energy	Incompatible requests on the same parameter
S2: radio-resource conflict	Power-control and resource-allocation	Power level, resource-block profile	Throughput, interference, loss	Different parameters influencing common KPMs
S3: mobility-energy conflict	Mobility-robustness and energy-saving	Handover offsets, sleep/wake	Handover failures, ping-pong, energy	Sleep decisions modify coverage and mobility
S4: load-balancing conflict	Load-balancing and mobility	Cell offsets, handover thresholds	Cell load, edge throughput, handovers	Joint control of user association
S5: multi-xApp stress	All xApps	Joint action space	All registered KPMs and SLAs	Mixed direct, indirect, and implicit relations

Table 5 Solver configurations per campaign environment. *Experiment processes* is the number of concurrent operating-system processes used to execute experimental runs; *num_search_workers* is the intra-solver search parallelism of each individual CP-SAT call. Both forms of parallelism are disabled throughout: every campaign executes as a single sequential process issuing single-threaded CP-SAT calls. Instance seeds are fixed and deterministic in all campaigns.

Campaign	OR-Tools	Experiment processes	num_search_workers	Instance seeds	Limit	Presolve
Static	9.14	1	1	42–91	100 ms	enabled
Dynamic	9.14	1	1	1000–	100 ms	enabled
Repair ablation	9.15	1	1	42–91	100 ms shared	enabled
Sequential	9.15	1	1	42–46	100 ms shared	enabled

core count of the host machine (28 on the primary Windows host, 2 on the secondary Linux host) affects neither instance identity nor search behaviour.

The static campaign nevertheless did not control garbage collection, CPU affinity, or background system load, so its latency results include host-level timing variability, this affects measured durations only, never the generated instances or the solver search path. Warm-up runs were executed prior to measurement to initialize the runtime environment.

Matched instances.

All methods within a static matched instance received the same domains, relations, utilities, and context: the harness generates each instance once per (scenario, seed) pair and evaluates every method against that identical instance before advancing. Host-level timing variability therefore affects measured latency but not instance identity, and the paired statistical analysis of Section 7 is well defined.

Latency boundaries.

Reported pipeline latency decomposes as

$$T_{\text{pipeline}} = T_{\text{update}} + T_{\text{prop}} + T_{\text{decode}} + T_{\text{repair}}, \quad (30)$$

where T_{update} covers instance update and model reconstruction, T_{prop} covers GAC propagation, T_{decode} covers candidate decoding, and T_{repair} covers repair-core expansion together with

$$T_{\text{repair}} = T_{\text{model-build}} + T_{\text{solve}} + T_{\text{extract}}. \quad (31)$$

Independent verification T_{verify} is measured separately and is *not* included in the reported figures, nor are certificate construction, serialization, or E2 transport. Accordingly, the reported quantities are *computational coordination latency, excluding certificate serialization and E2 transport*, and the 1.8 ms and 0.05–4.9 ms figures quoted throughout exclude verification. Where tables are labelled

“end-to-end”, the term denotes coverage of the complete coordination pipeline from instance update through repair, not a closed-loop control delay.

6.3 Structural characterization

The synthetic instances exhibit structural properties conducive to local repair. The constraint graph has a mean degree of 3.4 and maximum degree of 8, with relations having arities of 2 or 3. This sparse clustering may favour localized propagation and repair and therefore limits the generalizability of the sequential results.

6.4 Compared methods

All methods share the same registered input model, timing harness, and final verifier.

The evaluation includes four groups of methods.

1. **Sanity-check policies:** uncoordinated forwarding and static priority. These carry no verifier and establish the baseline rate at which unarbitrated proposals satisfy the registered model.
2. **Complete coordination baselines:** complete CP-SAT rebuilt from the full model, and GAC followed by complete CP-SAT.
3. **Repair variants:** fixed violated-scope, radius-based, frontier, explanation-guided, assumption-core, and full-scope repair. The default CONSISTXAPP configuration combines GAC, utility decoding, explanation-guided expanding repair, and verification.
4. **Alternative candidate and objective mechanisms:** continuous relaxation, GAC followed by utility decoding without repair, a QoS-inspired kept-request optimizer, a discretized Nash-social-welfare optimizer, and greedy, min-conflicts, and forward-checking candidate generators.

The QoS-aware and bargaining baselines reproduce objective structures inspired by previous work [4, 9], they are not reimplementations of the complete published systems.

In the sequential suite, the incremental candidate is the previous verified decision. It is re-verified under the updated model and repaired if necessary. The competing solver baselines rebuild and re-solve the full CP-SAT model, either without a hint or with the previous solution supplied as

a hint. Thus, the comparison is against the evaluated Python rebuild-and-resolve implementations, not against every possible persistent or incremental solver architecture.

6.5 Outcome terminology

Verified feasibility includes only actions accepted by the final verifier. A fallback is reported separately and is not silently counted as an optimized solution.

The 100 ms threshold is a post-hoc compliance metric checked after propagation, decoding, repair iterations, and verification. It is not a preemptive interrupt. Therefore, the experiments do not demonstrate hard real-time schedulability.

6.6 Metrics and statistical protocol

Filtering is measured using

$$r_D = 1 - \frac{\sum_i |D_i^*|}{\sum_i |D_i|} \quad (32)$$

and

$$r_R = 1 - \frac{\sum_c |R_c^*|}{\sum_c |R_c|}, \quad (33)$$

where R_c^* is the active part of relation R_c after propagation.

Binary paired outcomes are reported as counts and compared using exact McNemar tests. Continuous paired outcomes are summarized using medians, paired Wilcoxon signed-rank tests, and matched-pairs rank-biserial correlations. All statistical analyses were executed in the Python 3.10 Linux environment using `scipy.stats` (SciPy 1.11.0) and `statsmodels.stats.multitest` (Statsmodels 0.14.0), irrespective of the environment used to generate each campaign’s raw measurements. Holm correction is applied within the declared comparison families (Table 7). Statistical operations were implemented as follows:

- Wilcoxon signed-rank tests: `scipy.stats.wilcoxon` (SciPy 1.11.0).
- McNemar tests: `scipy.stats.binomtest` applied to discordant pairs (SciPy 1.11.0).
- BCa intervals: `scipy.stats.bootstrap` (SciPy 1.11.0).
- Holm adjustment: `multiplerepairs` from `statsmodels.stats.multitest` (Statsmodels 0.14.0).

Table 6 Decision-outcome terminology.

Term	Meaning
Direct decode	Candidate accepted without repair
Local repair	Candidate corrected using a proper subset of variables
Full-scope repair	Repair after expansion to the complete model
Fallback	Registered fallback accepted by the verifier
Blocked	No verifier-accepted action produced
Verified feasibility	Direct, local, full-scope, or verified fallback outcome
Deadline compliance	Verified action completed within the post-hoc threshold

- Rank-biserial correlation: custom exact formula implementation (provided in the artifact).
- Hodges–Lehmann estimate: custom Walsh-average implementation (provided in the artifact).

Bootstrap 95% confidence intervals for paired medians are computed using the paired-difference BCa resampling procedure with 10,000 replicates and seed 1. Marginal KPM confidence intervals are computed using a separate decision-level BCa bootstrap, also with 10,000 replicates.

Estimands and resampling units.

The phrase “paired median difference” is used throughout in the paired-design sense

$$\text{median}_i(A_i - B_i), \quad (34)$$

that is, the median of the per-unit differences, and never the difference of marginal medians $\text{median}(A_i) - \text{median}(B_i)$. The two quantities do not coincide in general.

The estimand and resampling unit differ by campaign:

- **Static comparisons:** the estimand is the median paired per-instance latency difference, the resampling unit is the generated instance.
- **Sequential comparisons:** the primary estimand is the paired difference between episode-level median per-epoch latencies; the resampling unit is the episode.
- **KPM intervals:** the estimand is a marginal median, the resampling unit is the decision. These intervals are descriptive.

In the sequential suite, the episode is the inferential unit for median-latency comparisons because

epochs from the same episode are temporally dependent. Epoch-level p95, p99, and maximum values are reported descriptively. With ten episodes of 30 epochs, especially the p99 estimate is based on very few extreme observations and should not be interpreted as a precise tail model.

Timeouts are treated as censored outcomes and are not converted into ordinary latency observations.

6.7 Execution environments

The scenario, dynamic, and temporal campaigns ran on the primary reference host: a 28-core Intel system under Windows 11, Python 3.13, and OR-Tools 9.14. The harness executed sequentially on this host; the core count is reported for host identification only and does not indicate parallel execution.

The operational-candidate ablation, cold-start, sequential, objective-gap, baseline, and verifier-validation campaigns ran in a 2-vCPU Linux container using Python 3.10 and OR-Tools 9.15.

All paired comparisons are within an environment. Absolute latency values from different environments are not directly compared. The artifact records the versions, seed hierarchy, and commands. A single-host rerun remains desirable before deployment conclusions are drawn.

7 Results

7.1 Trust layer

This subsection answers RQ1. All four verifier-validation gates passed with zero divergence (Table 8). The total comprised 60,650 exhaustive assignment comparisons, 1,800 triple-agreement checks, 602,000 property-based checks, and 16,000 oracle-labelled input mutants.

Table 7 Multiple-hypothesis correction families. Statistical significance claims are adjusted within these predefined families using the Holm–Bonferroni method. The p -values are reported for the primary confirmatory tests. For McNemar comparisons, b and c denote the off-diagonal discordant pairs: $b = \#(A \text{ succeeds}, B \text{ fails})$ and $c = \#(A \text{ fails}, B \text{ succeeds})$.

Family	Outcome	Method A	Method B	N	b	c	Raw p	Adjusted p
F1a	End-to-end verified feasibility	CP-SAT	CONSISTXAPP	250	0	0	1.000	1.000
	End-to-end verified feasibility	Continuous	CONSISTXAPP	250	0	7	0.016	0.031
F1b	Direct decoder feasibility	Min-conflicts	Utility decoder	250	177	0	$1.0 \cdot 10^{-53}$	$2.1 \cdot 10^{-53}$
	Direct decoder feasibility	Forward check	Utility decoder	250	213	0	$1.5 \cdot 10^{-64}$	$4.6 \cdot 10^{-64}$
	Direct decoder feasibility	Greedy	Utility decoder	250	14	2	$4.2 \cdot 10^{-3}$	$4.2 \cdot 10^{-3}$
F2	Repair success	Radius-zero	Explanation	248	0	21	$9.5 \cdot 10^{-7}$	$2.9 \cdot 10^{-6}$
	Repair success	Violated-scope	Explanation	248	0	21	$9.5 \cdot 10^{-7}$	$2.9 \cdot 10^{-6}$
	Repair success	Frontier	Explanation	248	0	0	1.000	1.000
F3	Repair latency	Full solve	Explanation	248	–	–	$8.0 \cdot 10^{-34}$	$8.0 \cdot 10^{-34}$
F4	Core size	Radius-zero	Explanation	248	–	–	$4.7 \cdot 10^{-5}$	$1.4 \cdot 10^{-4}$
	Core size	Violated-scope	Explanation	248	–	–	$4.7 \cdot 10^{-5}$	$1.4 \cdot 10^{-4}$
	Core size	Frontier	Explanation	248	–	–	$3.4 \cdot 10^{-3}$	$3.4 \cdot 10^{-3}$
F5	Seq. latency 2%/100	Re-solve	CONSISTXAPP	10	–	–	0.001953	0.007812
	Seq. latency 2%/300	Re-solve	CONSISTXAPP	10	–	–	0.001953	0.007812
	Seq. latency 2%/600	Re-solve	CONSISTXAPP	10	–	–	0.001953	0.007812
	Seq. latency 2%/1000	Re-solve	CONSISTXAPP	10	–	–	0.001953	0.007812

Table 8 Verifier-validation gates.

Gate	Content	Checks	Divergences
G1	Exhaustive oracle comparison	60,650	0
G2	Triple agreement with CP-SAT	1,800	0
G3	Property-based testing	602,000	0
G4	Oracle-labelled mutants	16,000	0
Total		680,450	0

On small instances for which exhaustive enumeration was practical, no value removed by propagation belonged to a feasible assignment. Incremental propagation reproduced the domains obtained by full recomputation on all 15,000 dynamic epochs. The explanation checker accepted 15,000 valid generated explanations (covering both wipeout and ordinary value-removal explanations) and rejected 15,000 synthetically corrupted counterparts (using random value-removal operators) with zero false accepts and zero false rejects. Operational audit certificates are implemented and replayable but their quantitative serialization overhead was not formally evaluated.

The scope of this implementation-level trust argument is discussed in Section 8.

7.2 Consistency filtering

Propagation removed 21.7% of candidate values and 34.1% of relation tuples when pooled over the static suite. Median GAC latency remained below 0.3 ms in every scenario.


7.3 Static-suite feasibility and latency

Table 11 reports integer decision outcomes. The conflict-unaware utility decoder directly produced a valid joint assignment on only 2 of 250 instances. CONSISTXAPP repaired the remaining 248 candidates and achieved 100% verified feasibility.

The two unarbitrated sanity-check policies are reported separately in Table 10, because they carry no verifier and therefore cannot produce a pipeline

Table 9 Consistency-propagation results over 50 generated instances per scenario.

Scenario	Values before	Values after	r_D (%)	Tuples before	Tuples after	r_R (%)	Median (ms)	P95 (ms)
S1	1800	1414	21.4	2403	1713	28.7	0.07	0.09
S2	3600	2676	25.7	6083	3697	39.2	0.15	0.18
S3	4000	3148	21.3	7419	4858	34.5	0.17	0.20
S4	4800	3908	18.6	9694	6846	29.4	0.20	0.24
S5	6000	4662	22.3	11747	7488	36.3	0.26	0.31

**Fig. 2** Domain and relation-table reductions obtained by generalized arc consistency.**Table 10** Sanity-check policies without arbitration. The reported quantity is the number of raw policy outputs that satisfy the registered hard model, not a pipeline outcome.

Raw policy	Registered-model feasible outputs
Uncoordinated	0/250
Static priority	0/250

outcome in the sense of Table 11. What they measure is the rate at which raw proposals happen to satisfy the registered model: zero of 250 in both cases. Under the common forwarding gate, such outputs are not admitted, since the gate admits only verifier-accepted candidates.

Complete CP-SAT and GAC followed by full CP-SAT also achieved 100% verified feasibility. Their median latencies were 2.8 and 2.6 ms, respectively, compared with 1.8 ms for CONSISTXAPP.

The paired median difference between CONSISTXAPP and complete CP-SAT was -0.85 ms, with a bootstrap 95% interval $[-0.93, -0.72]$ ms. The paired Wilcoxon result was $p \approx 7 \cdot 10^{-34}$ and the matched-pairs rank-biserial correlation

was -0.89 . This difference is statistically significant under the specified paired test, although operationally modest in magnitude.

The objective-inspired baselines also achieved 100% verified feasibility and retained approximately 64% of preferred actions. Their latency results were obtained in the secondary environment and are not directly compared with results from the primary host.

The continuous-relaxation stage was counter-productive on this suite. It increased median latency to approximately 80 ms, reduced post-hoc deadline compliance to approximately 57%, and left six or seven instances blocked, depending on the core rule. Its feasibility deficit relative to the discrete pipeline was statistically significant (exact McNemar $p = 0.016$ unadjusted, $p_{\text{Holm}} = 0.031$). It was therefore excluded from the proposed operational pipeline.

For the continuous-relaxation evaluations, the upstream end-to-end coordination feasibility sample size is strictly $N = 250$. However, the downstream KPM-analysis sample is $N = 249$ because one verifier-accepted decision failed during the post-verification KPM simulation stage due to a simulator numerical error. A post-verification simulator error does not constitute a failure of the upstream coordination pipeline. The failed simulation concerns a single S5 instance that was verifier-accepted under both CONSISTXAPP and the continuous variants, the CP-SAT KPM simulation on the same generated instance completed, so its marginal sample remains $n = 250$, whereas the CONSISTXAPP sample excludes that instance ($n = 249$) and the explanation-guided continuous sample excludes it as well ($n = 242$ of its 243 verifier-accepted decisions). Missing simulator values were omitted pairwise in computations using decisions as the resampling unit, and all KPM intervals are marginal decision-level bootstrap intervals rather than paired-sample intervals.

Table 11 Coordination-pipeline outcomes on the static suite, covering instance update through repair and excluding independent verification, certificate serialization, and E2 transport. Deadline compliance is post-hoc. Sanity-check policies are reported separately in Table 10.

Method	N	Direct	Local repair	Full scope	Fallback	Blocked	Feasible (%)	≤ 100 ms (%)	Median (ms)	P95 (ms)
GAC + utility decode	250	2	0	0	0	248	0.8	—	—	—
CP-SAT	250	0	0	250	0	0	100.0	100.0	2.8	4.8
GAC + CP-SAT	250	2	0	248	0	0	100.0	100.0	2.6	4.7
ConsistXApp	250	2	248	0	0	0	100.0	100.0	1.8	3.8
Continuous only	250	189	0	0	0	61	75.6	23.6	180.4	365.7
GAC + continuous decode	250	186	0	0	0	64	74.4	44.4	80.6	251.3
Continuous + repair (explanation)	250	186	57	0	0	7	97.2	57.4	79.9	253.3
Continuous + repair (radius)	250	186	57	0	0	7	97.2	57.0	82.6	251.9
Continuous + repair (assumption)	250	186	58	0	0	6	97.6	57.8	81.3	254.4

Table 12 Final certification status of completed lexicographic repair sequences. A decision or epoch is LEX-OPTIMAL when every lexicographic level of its repair returned OPTIMAL; FEASIBLE-NOT-CERTIFIED denotes a verifier-accepted incumbent returned before full lexicographic certification. An episode is counted as LEX-OPTIMAL only if every repaired epoch in that episode achieved complete lexicographic certification. Counts are completed repair episodes, not internal solver calls, each sequence comprises one or more core-feasibility calls followed by up to three objective-level optimizations.

Campaign	Inferential unit	Completed repair sequences	Fully lex- optimal	Feasible-not- certified
Static	decision	248	248	0
Dynamic	epoch	14,755	14,755	0
Sequential F5 subset	episode	40	40	0

7.4 Repair-core ablation

Table 14 compares repair-core strategies on all 248 repair candidates produced by the conflict-unaware utility decoder on the static suite. Fixed violated-scope and radius-zero repair succeeded on 227 of 248 cases (91.5%), whereas the expanding variants succeeded on all 248. The exact paired McNemar comparison between explanation-guided and radius-zero repair contained 21 discordant pairs, all favouring explanation-guided repair ($b = 0$, $c = 21$), giving $p_{\text{raw}} = 9.5 \cdot 10^{-7}$ and $p_{\text{Holm}} = 2.9 \cdot 10^{-6}$ within family F2. The success-rate advantage of core expansion over fixed-scope repair is therefore statistically significant on this suite. This advantage is attributable to expansion itself rather than to explanation guidance specifically: the frontier, assumption-core, and full-scope strategies also resolved all 248 candidates. The specific benefit of explanation guidance is that it recovered the 21 missed candidates with markedly smaller final cores than full-scope solving.

Explanation-guided expansion used a median final core of 11 variables, compared with 16 for radius two and 20 for full-scope solving. Full-scope repair required a median 1.78 ms, compared with 1.19 ms for explanation-guided repair. The




Fig. 3 Final repair-core sizes across the 248 operational-candidate repairs.

latency advantage of explanation-guided repair over full-scope solving is statistically significant ($p = 8.0 \cdot 10^{-34}$, matched-pairs rank-biserial correlation -0.89 , family F3 in Table 7). Its final cores are moderately but significantly larger than the fixed violated-scope and radius-zero cores (family F4: $p_{\text{Holm}} = 1.4 \cdot 10^{-4}$), which is the price of the expansions that recovered the 21 cases lost by fixed-scope repair.

Table 13 Paired static-instance comparisons. Latency effect sizes are matched-pairs rank-biserial correlations. HL denotes the Hodges–Lehmann paired difference in milliseconds. When $b + c = 0$ no discordant pairs exist, the McNemar comparison is reported as $p = 1$ by convention and carries no evidence of superiority in either direction.

Comparison	N	Feasibility A/B	b	c	p_{McN}	Median A/B	HL	p_{Wilc}	r_{rb}
CONSISTXAPP vs CP-SAT	250	100/100	0	0	1.000	1.76/2.76	−0.85	$< 10^{-33}$	−0.89
CONSISTXAPP vs GAC +CP-SAT	250	100/100	0	0	–	1.76/2.63	−0.75	$< 10^{-30}$	−0.85
Continuous+repair vs CONSISTXAPP	250	97.2/100	0	7	0.031	79.9/1.75	+94	$< 10^{-40}$	+1.0
Continuous+repair vs CP-SAT	250	97.2/100	0	7	0.031	79.9/2.76	+93	$< 10^{-40}$	+1.0

Table 14 Repair-core ablation. “Used” gives the number of expansions driven by explanations (e), assumption cores (a), and the frontier rule (f).

Strategy	Operational candidates (248)					Stress suite (50)		
	Success (%)	Final core	Used (e/a/f)	Released	Median (ms)	Success (%)	Final core	Median (ms)
Violated	91.5	10	0/0/0	0.00	0.97	100	8	0.54
Radius 0	91.5	10	0/0/0	0.00	1.01	100	8	0.51
Radius 1	99.2	14.5	0/0/0	0.00	1.57	100	16	0.65
Radius 2	100.0	16	0/0/0	0.00	1.71	100	22	0.90
Frontier	100.0	11	0/0/21	0.48	1.23	100	8	0.64
Explanation	100.0	11	16/0/8	0.30	1.19	100	8	0.64
Assumption	100.0	10.5	0/21/0	0.15	2.44	100	8	2.50
Full scope	100.0	20	0/0/0	0.00	1.78	100	24	1.37

These results concern candidates produced by the deliberately conflict-unaware decoder. They demonstrate the effect of expansion relative to fixed-boundary repair, they do not show that 248 repairs would be required under a stronger operational decoder.

7.5 Objective quality

The local repair boundary can exclude globally better solutions. Across the static scenario suite, explanation-guided repair retained 53.8% of preferred requests on average, whereas full objective optimization retained 64.6%. The mean difference was therefore 10.8 percentage points. The paired asymptotic Wilcoxon result was $p < 10^{-35}$. Because 12.8% of paired differences were zero out of the $N = 250$ static instances, the effective non-zero sample was $N_{\text{eff}} = 218$, we used the asymptotic Wilcoxon signed-rank test with `zero_method='wilcox'`. The median paired gap was 11.2 percentage points (IQR: 6.8–14.5 percentage points) in favour of full optimization, with a matched-pairs rank-biserial correlation of -0.92 . The gap did not exceed 40 percentage points. Median latency was 1.25 ms for local repair and

2.21 ms for the full objective optimizer. This 1.25 ms figure is not directly comparable with the 1.19 ms reported in Section 7.4: the 1.19 ms result concerns the repair-core strategy ablation, which measures core feasibility only, whereas the 1.25 ms result belongs to the separate objective-quality campaign and additionally includes the operational lexicographic optimization stages.

Accordingly, the two capabilities should not be conflated. CONSISTXAPP provides verified feasible repair with an audit trace. A complete objective solve provides global optimality when the solver returns the required optimal status for the complete encoded objective.

7.6 Incremental propagation

Incremental propagation reproduced the fully recomputed filtered domains on all 15,000 dynamic epochs.

The largest relative reductions occurred for restrictive and utility-only transitions. Relaxing and mixed transitions gained less because the affected component was first restored. In absolute terms, the savings were fractions of a millisecond. The larger sequential benefit reported next

Table 15 Incremental propagation versus full recomputation. Times are mean propagation times per epoch.

Transition	Count	Recomputed (ms)	Incremental (ms)	Gain (%)	Restored	Agreement (%)
Base	150	–	–	–	–	100.0
Restrictive	750	0.102	0.014	86.3	0.00	100.0
Relaxing	750	0.100	0.050	50.2	0.38	100.0
Mixed	12,600	0.087	0.059	32.0	0.67	100.0
Utility-only	750	0.099	0.004	95.9	0.00	100.0

is therefore mainly due to avoiding complete model reconstruction and re-solving, rather than propagation alone.

7.7 Conflict-aware decoders

The utility decoder is a weak candidate generator. Greedy checking, min-conflicts, and forward checking were therefore evaluated on the same 250 static instances.

Min-conflicts and forward checking produced directly feasible candidates on 71.6% and 86.0% of instances, respectively. Thus, the 248 repairs emphasized in the utility-decoder ablation should not be read as the repair rate of the strongest evaluated candidate generator.


Neither conflict-aware decoder eliminated repair. Forward checking left 35 candidates unresolved within its 50 ms budget, and direct feasibility fell from 100% on S1 to 62% on S5. Conflict-aware decoders also retained fewer preferred values than the utility decoder.

Explanation-guided repair completed the candidates from all four decoders to 100% verified feasibility in this suite. A practical pipeline should therefore combine a conflict-aware candidate generator with verified repair (see Section 8 for the decoder-dependence caveat).

7.8 Cold-start and sequential operation

Figure 4 reports the cold-start results over twenty seeds per size. On isolated near-threshold instances, complete CP-SAT remained within the 100 ms post-hoc threshold up to $n = 1000$, with a median latency of 38 ms and a p95 of 43 ms. The one-shot CONSISTXAPP pipeline required a median of 68 ms and a p95 of 90 ms. Its observable Python GAC stage accounted for approximately 40% of the measured cost.

Thus, the evaluated pipeline provides no cold-start latency advantage over complete CP-SAT for

**Fig. 4** Cold-start latency on generated near-threshold instances. Lines show medians and shaded regions extend to the pooled p95 over twenty seeds per size. On these isolated instances, the one-shot CONSISTXAPP pipeline offers no latency advantage over complete CP-SAT.

these isolated instances. A deployment that faces unrelated coordination problems should use the complete solver unless it requires another feature of the pipeline, such as its verification or audit trace.

The sequential experiment evaluates a different operating regime. Consecutive epochs retain the same variables, domains, and utilities, while a fraction of the registered relations changes. At each epoch, the incremental strategy first re-verifies the previous verified decision. If that decision violates the updated model, it is repaired. The two comparison methods rebuild and solve the complete CP-SAT model, either without a hint or with the preceding solution supplied as a hint.

The experiment is not a trivial re-verification workload. At $n = 1000$ and $\delta = 2\%$, the previous decision remained feasible without repair on only 1.7% of epochs. At $n = 100$, the corresponding fraction was 69%.

For $\delta = 2\%$, verified incremental repair required a median of 0.05 ms per epoch at $n = 100$ and 4.9 ms at $n = 1000$. Rebuilding and re-solving

Table 16 Conflict-aware candidate generators. “Direct” is verifier-accepted feasibility before repair. “Repair needed” is the number of candidates subsequently passed to explanation-guided repair.

Decoder	Direct (%)	Kept preference (%)	Median (ms)	P95 (ms)	Repair needed	Median core
Utility decode	0.8	77.5	0.15	0.29	248	11
Greedy + checks	5.6	71.2	0.05	0.07	236	8
Min-conflicts	71.6	55.1	0.37	4.50	71	4
Forward checking	86.0	57.3	0.19	50.2	35	12

the complete model required medians from 2.8 to 25.6 ms. The hinted rebuild-and-resolve baseline required medians from 2.9 to 27.4 ms.

All three methods reached 100% verified feasibility on the evaluated sequential epochs. Treating the episode as the inferential unit, the per-episode median favoured incremental repair in all ten paired episodes at each evaluated size and change rate. With ten non-zero paired differences in the same direction, the two-sided exact Wilcoxon signed-rank result attained its minimum possible value ($p = 0.001953$) for ten pairs. Table 17 lists the resulting maximum Holm-adjusted p -values across all four evaluated configurations in comparison family F5. The exact tests handled zero differences by omitting them from the ranks, and no ties were present. The matched-pairs rank-biserial correlation was -1.0 .

At $n = 1000$ and $\delta = 2\%$, the Hodges–Lehmann estimate of the per-epoch latency difference between incremental repair and re-solving was -20.3 ms.

At $n = 1000$, the unhinted re-solve spent a median of 6.9 ms constructing the Python model and 18.4 ms in solver search. In this implementation, both components scaled with complete model size, and a solution hint removed neither model reconstruction nor complete model processing (see Section 8 for the scope of this baseline).

The pooled tail measurements reveal an important limitation. At $n = 1000$ and $\delta = 2\%$, the p95 values of incremental repair and unhinted re-solving were 30.6 and 31.2 ms, respectively. The pooled p99 values were approximately 84 and 33 ms, and the largest observed incremental-repair latency was 129 ms. At $\delta = 10\%$, the pooled incremental-repair p95 increased to approximately 60 ms.

These observations suggest that core growth can remove the median advantage on difficult




Fig. 5 Per-epoch latency when 2% of the relations change between epochs. Lines show pooled medians and shaded regions extend to the pooled p95. The methods use identical generated episodes. The figure describes the evaluated Python implementations and is not a hard real-time guarantee.

epochs, the small number of extreme observations underlying the pooled p99 is discussed in Section 8.

The drift variant reduces the risk that a fixed planted assignment makes sequential repair artificially easy. Re-planting 1% of the variables invalidated the previous decision on every epoch. At $n = 1000$, incremental repair still had a lower median latency, 9.8 ms compared with 28.1 ms for unhinted re-solving. Its pooled p95 and p99 were nevertheless 43.8 and 81.3 ms, preserving the same tail qualification.

Separately from the operational lexicographic pipeline, a diagnostic temporal-stability sweep evaluated the scalarized formulation $U - \mu H$ of (7) on the dynamic sequence. Increasing the churn weight μ from zero to one reduced mean action churn by 11%, while measured throughput changed by less than 0.2% and registered-model feasibility was preserved. The sweep therefore characterizes the scalarized diagnostic model, not the lexicographic

Table 17 Exact sequential Wilcoxon signed-rank results (Family F5). The Holm correction adjusts the raw exact $p = 0.001953$ for four simultaneous comparisons.

Size	Comparison	Episodes	Raw p	Holm p	r_{rb}
2% / 100	Incremental vs Re-solve	10	0.001953	0.007812	-1.0
2% / 300	Incremental vs Re-solve	10	0.001953	0.007812	-1.0
2% / 600	Incremental vs Re-solve	10	0.001953	0.007812	-1.0
2% / 1000	Incremental vs Re-solve	10	0.001953	0.007812	-1.0

Table 18 Per-epoch latency over ten paired episodes of thirty epochs. Entries are pooled median and p95 in milliseconds, the value after the semicolon for incremental repair is the pooled p99. All strategies reached 100% verified feasibility. Because epochs within an episode are dependent and only 300 epochs are available per configuration, p95 and especially p99 are descriptive rather than precise tail estimates. Drift rows re-plant 1% of variables per epoch.

Change / size	CP-SAT re-solve	CP-SAT + hints	Incremental repair
2% / 100	2.8 (4.8)	2.9 (5.2)	0.05 (1.8, 4.4)
2% / 300	8.2 (11.4)	8.7 (11.8)	1.4 (5.0, 13.2)
2% / 600	15.2 (19.5)	16.3 (20.7)	2.7 (10.9, 22.7)
2% / 1000	25.6 (31.2)	27.4 (32.8)	4.9 (30.6, 83.7)
10% / 1000	18.3 (27.1)	20.2 (28.6)	4.5 (59.6, 137.8)
Drift / 600	17.8 (20.9)	18.9 (22.0)	4.0 (25.8, 69.3)
Drift / 1000	28.1 (32.5)	29.8 (34.2)	9.8 (43.8, 81.3)

operational objective, whose ordering is invariant to positive scaling of the primary level. This is a result of the synthetic pseudo-physical model and should not be interpreted as a physical-network performance guarantee.

7.9 Network-level descriptive results

Table 19 reports descriptive KPM estimates for verifier-accepted decisions. These values were generated by the pseudo-physical model. They are not measurements from an emulated or deployed O-RAN system.

The marginal bootstrap intervals overlap for some KPMs, but overlap of separate confidence intervals is not a paired equivalence test and does not establish absence of a material difference. No non-inferiority or equivalence margin was defined before analysis. Consequently, Table 19 supports only a descriptive comparison.

The principal experimentally supported distinction is registered-model feasibility, not KPM superiority, the additional evidence needed for KPM equivalence claims is stated in Section 8.

8 Limitations and Threats to Validity

This section is the single authoritative statement of the validity caveats of this study, earlier sections refer here rather than restating them.

8.1 Registered-model validity

Every formal guarantee in this paper is conditional on the registered model. A verified decision satisfies the registered domains, relations, model version, and context. It is not thereby safe for a physical RAN if the model omits a relevant dependency or uses an invalid threshold.

This limitation is particularly important for implicit conflicts, whose representation depends on the quality of a parameter-KPM model, validated profile, digital twin, simulator, or conservative operator rule. The independent verifier checks consistency with the registered representation, it does not validate the representation against physical reality.

Model-version binding is also required to prevent a time-of-check/time-of-use inconsistency. A candidate verified against one model version must not be forwarded after the registered model has changed unless it is re-verified.

Table 19 Descriptive KPMs of verifier-accepted decisions, presented as preliminary diagnostic evidence, reported as mean and bootstrap 95% interval. Energy is measured in $\mu\text{J}/\text{bit}$. The continuous-repair sample contains 242 verified decisions. One S5 CONSISTXAPP KPM evaluation failed numerically after verification, the CP-SAT KPM for that instance remained available. The raw uncoordinated proposal had a counterfactual mean throughput of 11.7 Mbps, but no uncoordinated proposal was accepted for execution.

Method	Throughput (Mbps)	Energy ($\mu\text{J}/\text{bit}$)	Jain fairness
CP-SAT ($n = 250$)	10.84 [10.56, 11.12]	50.8 [49.6, 52.1]	0.443 [0.433, 0.452]
CONSISTXAPP ($n = 249$)	10.41 [10.15, 10.68]	52.2 [51.0, 53.5]	0.433 [0.424, 0.443]
Continuous relaxation + explanation-guided repair ($n = 242$)	10.93 [10.68, 11.18]	50.6 [49.3, 51.9]	0.451 [0.441, 0.461]

8.2 Verifier evidence

The verifier-validation gates provide substantial implementation evidence but do not constitute a proof of universal correctness. The oracle, production verifier, and declarative implementation are independent at the parser and checking levels, but the generated test distribution still limits the conclusions.

Input mutation tests the classification of modified inputs. It is not code-mutation testing and does not measure how many implementation faults the test suite would detect. Zero divergence over 680,450 checks therefore means zero observed disagreement on the tested inputs, not zero probability of an undiscovered defect.

8.3 Local consistency and repair completeness

GAC is a local property. Non-empty filtered domains do not establish global feasibility. Similarly, infeasibility of a proper-core repair model may be caused by fixed boundary values and does not establish infeasibility of the complete registered model.

Conditional completeness applies only after expansion to the complete variable set, exact encoding of the registered finite-domain relations, and a conclusive solver status before the time limit. A solver status of **UNKNOWN** must not be interpreted as infeasibility.

The operational audit certificates establish replayable feasibility and expansion traces. They do not establish minimum repair-core size or global objective optimality. Certificate serialization size and replay latency were not separately measured.

8.4 Objective encoding and solver status

The operational repair stage uses a lexicographic preference order, but feasibility and optimality must be distinguished. A **FEASIBLE** CP-SAT result establishes a feasible assignment, whereas an **OPTIMAL** result is required to establish optimality for the encoded objective.

The separate objective-gap experiment compares repair with a complete objective optimizer. Any final archival artifact must record the integer scaling, coefficient bounds, lexicographic encoding, solver status, time limit, number of search workers, random seed, and presolve settings used in that experiment. Without these details, the reproducibility package is incomplete.

8.5 Decoder dependence

The utility decoder was deliberately conflict-unaware. It generated 248 repair cases out of 250 instances, whereas forward checking reduced the number requiring repair to 35. Therefore, the headline repair count measures the contribution of repair under that weak decoder, not the repair frequency of the strongest candidate generator.

The conflict-aware decoder experiment shows that repair remains useful, but the large-scale sequential campaign was not repeated with every decoder. A production evaluation should combine a conflict-aware candidate generator with verified repair and should report end-to-end latency, repair frequency, objective quality, certificate size, and tail behaviour for that complete pipeline.

8.6 Generated-instance structure

The scale experiments use planted-feasible positive-table instances with fixed domain size, relation

count, and tightness. These choices make coordination failures distinguishable from model infeasibility, but they may favour local repair.

The experiments do not systematically vary graph density, clustering, relation arity, domain size, tightness, feasible-region fragmentation, or the distance between consecutive feasible regions. They also do not evaluate large-scale sequences in which domains, xApp activation, utilities, and relations all change simultaneously.

The sequential conclusion must therefore be restricted to the evaluated generated workloads. It does not establish the same scaling behaviour for dense, high-arity, adversarial, or commercial multi-vendor O-RAN models.

8.7 Baseline scope and latency measurement

This subsection is the authoritative treatment of baseline scope; earlier mentions are deliberately brief and defer here.

The sequential baselines rebuild their Python CP-SAT models at every epoch. The hinted variant adds the previous solution as a hint but does not preserve a persistent solver state. Consequently, the experiment compares CONSISTXAPP with rebuild-and-resolve CP-SAT, with and without hints. It does not rule out faster persistent, incremental, or lower-level solver implementations. We do not claim that no persistent approach exists: we claim only that the evaluated Python rebuild-and-resolve implementations behave as reported. A solver-native large-neighbourhood baseline centred on the previous decision, and a persistent model supporting in-place relation mutation, are the two most informative comparisons not carried out here.

Reported latencies follow the decomposition of Equation (30) and exclude independent verification, certificate serialization, and E2 transport. They are therefore computational coordination latencies rather than closed-loop control delays.

The campaigns also ran in two documented execution environments. Paired comparisons remain within an environment, but absolute latencies across environments are not directly comparable. A single-host replication with fixed CPU affinity, warm-up, garbage-collection policy, and repeated timing runs would strengthen the systems conclusions.

8.8 Tail latency and real-time interpretation

The median sequential advantage is supported at the episode level. The p95, p99, and maximum results are descriptive pooled-epoch statistics. In particular, the p99 estimate for a 300-epoch configuration is determined by approximately three observations and has no reported dependence-aware uncertainty interval.

The 100 ms threshold is post-hoc. A stage can exceed the threshold before the overrun is detected. Hence, the current implementation does not provide a hard real-time guarantee, even though all verified decisions in several campaigns completed within 100 ms.

A production architecture requires a preemptive budget policy, for example per-stage time reservations, a solver limit based on remaining budget, an already verified fallback, or an escalation rule triggered by repair-core growth. No concrete tail-bounding policy was evaluated in this study.

8.9 Simulation and O-RAN deployment validity

The pseudo-physical model does not capture E2 transport delay, Near-RT RIC scheduling, container interference, message loss, hardware acceleration, or vendor-specific O-CU/O-DU/O-RU behaviour. The reported latencies are computational pipeline measurements, not complete closed-loop control delays.

No claim of physical-network safety, KPM equivalence, or deployment readiness follows from the current experiments. Closed-loop validation on an O-RAN emulator or experimental platform is required.

9 Conclusion

This paper formulated registered multi-xApp arbitration in the Near-RT RIC as a dynamic finite-domain constraint problem and presented CONSISTXAPP, a pipeline whose core contribution is an independently validated forwarding verifier with replayable audit certificates, supported upstream by incremental generalized arc consistency and explanation-guided repair. The stated formal properties - soundness of the forwarding gate relative to the registered model, solution preservation under

filtering, and conditional completeness of expanding repair cores - are standard, but they delimit precisely what the architecture does and does not guarantee.

The production verifier showed zero divergence over 680,450 checks against multiple independent implementations and oracles, providing substantial implementation-level agreement on the tested generated inputs. On 250 generated static instances in a single-worker deterministic environment (`num_search_workers = 1`), CONSISTXAPP reached a median computational latency of 1.8 ms. Thus, CONSISTXAPP outperformed the evaluated Python rebuild-and-resolve CP-SAT implementations on the generated sequential workloads. The operational pipeline repaired all 248 utility-decoder candidates. Across these repairs, explanation-guided expansion used markedly smaller cores than full-scope solving and succeeded on 21 candidates where fixed-scope repair failed. Stronger conflict-aware candidate generators substantially reduced the need for repair altogether.

The principal observed advantage was sequential and amortized. When 2% of relations changed between generated epochs, incremental repair required median latencies of 0.05–4.9 ms for 100–1,000 variables, compared with 2.8–25.6 ms for rebuilding and re-solving. This median advantage was consistent across paired episodes, demonstrating the efficiency of exploiting temporal stability when successive control requests largely preserve the previous constraints and assignments. At 1,000 variables with 2% relation changes, however, incremental-repair latency approached re-solve latency near the pooled p95 and exceeded it at the descriptive pooled p99.

However, several operational boundaries remain. Cold-start instances are better served by complete solving, local repair necessarily trades objective quality for speed (retaining fewer preferred requests than global optimization), and the observed median latency advantages do not constitute hard real-time tail guarantees. Future work will integrate conflict-aware decoding and verified repair as a unified pipeline, measure serialization overheads for audit certificates, evaluate denser constraint models, and perform closed-loop validation on O-RAN emulators and experimental testbeds.

Statements and Declarations

Funding

No funding was received for conducting this study.

Competing interests

The authors declare that they have no competing financial or non-financial interests that are directly or indirectly related to this work.

Author contributions

Lillia Ouali: Conceptualization, Validation, Supervision. **Madani Bezoui:** Conceptualization, Methodology, Formal analysis, Software, Investigation, Validation, Visualization, Writing – original draft, Writing – review & editing. **Kamal Amroun:** Validation, Writing – review & editing. **Ahcene Bounceur:** Validation, Supervision. All authors reviewed and approved the manuscript.

Data and code availability

The code, generated instances, raw logs, explanation traces, verifier-validation reports, reproducibility manifests, seed hierarchy, command lines, and scripts reproducing the tables, figures, and statistical tests are available in the public repository: <https://github.com/MadBezoui/OpenRanXAPP>.

The review artifact records the exact solver configuration for every campaign, including time limits, worker counts, random seeds, presolve settings, objective scaling, and solver statuses, in its reproducibility manifests.

Ethics approval

Not applicable.

Consent to participate

Not applicable.

Consent for publication

Not applicable.

Use of generative artificial intelligence

Generative language models were used to assist with manuscript organization, language editing,

and review of analysis scripts. The authors reviewed and corrected the resulting material and assume full responsibility for the scientific content, formal statements, references, experimental protocol, numerical results, and submitted version.

References

- [1] Polese, M., Bonati, L., D'Oro, S., Basagni, S., Melodia, T.: Understanding O-RAN: Architecture, interfaces, algorithms, security, and research challenges. *IEEE Communications Surveys & Tutorials* **25**(2), 1376–1411 (2023) <https://doi.org/10.1109/COMST.2023.3239220>
- [2] O-RAN Alliance: O-RAN Conflict Mitigation, version 1.00. Technical report, O-RAN Alliance, Near-Real-Time RIC and E2 Interface Work Group (2024). Specification portal accessed 14 July 2024. <https://specifications.o-ran.org/>
- [3] Adamczyk, C., Kliks, A.: Conflict mitigation framework and conflict detection in O-RAN Near-RT RIC. *IEEE Communications Magazine* **61**(12), 199–205 (2023) <https://doi.org/10.1109/MCOM.018.2200752>
- [4] Wadud, A., Golpayegani, F., Afraz, N.: QACM: QoS-aware xApp conflict mitigation in Open RAN. *IEEE Transactions on Green Communications and Networking* **8**(3), 978–993 (2024) <https://doi.org/10.1109/TGCN.2024.3431945>
- [5] Polese, M., Bonati, L., D'Oro, S., Basagni, S., Melodia, T.: ColO-RAN: Developing machine learning-based xapps for open ran closed-loop control on programmable experimental platforms. *IEEE Transactions on Mobile Computing* **22**(10), 5787–5800 (2023) <https://doi.org/10.1109/TMC.2022.3188013>
- [6] Bonati, L., Polese, M., D'Oro, S., Basagni, S., Melodia, T.: OpenRAN Gym: AI/ML development, data collection, and testing for O-RAN on PAWR platforms. *Computer Networks* **220**, 109502 (2023) <https://doi.org/10.1016/j.comnet.2022.109502>
- [7] Schmidt, R., Irazabal, M., Nikaein, N.: FlexRIC: An SDK for next-generation SD-RANs. In: *Proceedings of the 17th International Conference on Emerging Networking Experiments and Technologies. CoNEXT '21*, pp. 411–425. Association for Computing Machinery, New York, NY, USA (2021). <https://doi.org/10.1145/3485983.3494870>
- [8] Corici, M., Modroiu, R., Eichhorn, F., Troudt, E., Magedanz, T.: Towards efficient conflict mitigation in the converged 6G Open RAN control plane. *Annals of Telecommunications* **79**(9–10), 621–631 (2024) <https://doi.org/10.1007/s12243-024-01036-2>
- [9] Wadud, A., Golpayegani, F., Afraz, N.: Conflict management in the Near-RT-RIC of Open RAN: A game-theoretic approach. In: *2023 IEEE International Conferences on Internet of Things (iThings) and Related Conferences*, pp. 479–486. IEEE, Danzhou, China (2023). <https://doi.org/10.1109/iThings-GreenCom-CPSCom-SmartData-Cybermatics60724.2023.00095>
- [10] Brach del Prever, P., D'Oro, S., Bonati, L., Polese, M., Tsampazi, M., Lehmann, H., Melodia, T.: PACIFISTA: Conflict evaluation and management in Open RAN. *IEEE Transactions on Mobile Computing* **24**(10), 10590–10605 (2025) <https://doi.org/10.1109/TMC.2025.3570632>
- [11] Giannopoulos, A.E., Spantideas, S.T., Levis, G., Kalafatelis, A.S., Trakadas, P.: COMIX: Generalized conflict management in O-RAN xapps—architecture, workflow, and a power control case. *IEEE Access* **13**, 116684–116700 (2025) <https://doi.org/10.1109/ACCESS.2025.3585774>
- [12] Wadud, A., Golpayegani, F., Afraz, N.: xApp-level conflict mitigation in O-RAN: A mobility-driven energy-saving case. In: *IEEE INFOCOM 2025 – IEEE Conference on Computer Communications Workshops (INFOCOM WKSHPS)*. IEEE, London, UK (2025). <https://doi.org/10.1109/INFOCOMWKSHPS65812.2025.11153006>
- [13] Varshney, N., Puligheddu, C., Badawy, A.,

- Chiasserini, C.F.: Explainable artificial intelligence for conflict management: XAI4C for conflict detection and mitigation in O-RAN Near-RT RIC. *IEEE Vehicular Technology Magazine* **20**(4), 46–55 (2025) <https://doi.org/10.1109/MVT.2025.3625693>
- [14] Yap, R.H.C., Xia, W., Wang, R.: Generalized arc consistency algorithms for table constraints: A summary of algorithmic ideas. In: *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, pp. 13590–13597 (2020). <https://doi.org/10.1609/aaai.v34i09.7086>
- [15] Lecoutre, C.: STR2: Optimized simple tabular reduction for table constraints. *Constraints* **16**(4), 341–371 (2011) <https://doi.org/10.1007/s10601-011-9107-6>
- [16] Bessière, C.: Arc-consistency in dynamic constraint satisfaction problems. In: *Proceedings of the Ninth National Conference on Artificial Intelligence*, pp. 221–226. AAAI Press, Anaheim, CA, USA (1991)
- [17] Berlandier, P., Neveu, B.: Maintaining arc consistency through constraint retraction. In: *Proceedings of the Sixth IEEE International Conference on Tools with Artificial Intelligence*, pp. 426–431. IEEE, New Orleans, LA, USA (1994)
- [18] Debruyne, R.: Arc-consistency in dynamic CSPs is no more prohibitive. In: *Proceedings of the Eighth IEEE International Conference on Tools with Artificial Intelligence*, pp. 299–306. IEEE, Toulouse, France (1996)
- [19] Dev Gupta, S., Genc, B., O’Sullivan, B.: Explanation in constraint satisfaction: A survey. In: *Proceedings of the Thirtieth International Joint Conference on Artificial Intelligence*, pp. 4400–4407. International Joint Conferences on Artificial Intelligence Organization, Montreal, Canada (2021). <https://doi.org/10.24963/ijcai.2021/601>
- [20] Ginsberg, M.L.: Dynamic backtracking. *Journal of Artificial Intelligence Research* **1**, 25–46 (1993)
- [21] Shaw, P.: Using constraint programming and local search methods to solve vehicle routing problems. In: Maher, M., Puget, J.-F. (eds.) *Principles and Practice of Constraint Programming — CP98. Lecture Notes in Computer Science*, vol. 1520, pp. 417–431. Springer, Berlin, Heidelberg (1998). https://doi.org/10.1007/3-540-49481-2_30
- [22] Perron, L., Didier, F., Gay, S.: The CP-SAT-LP solver. In: Schaus, P. (ed.) *29th International Conference on Principles and Practice of Constraint Programming. Leibniz International Proceedings in Informatics*, vol. 280. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, Dagstuhl, Germany (2023). <https://doi.org/10.4230/LIPIcs.CP.2023.3>
- [23] Google OR-Tools Team: CP-SAT Solver. Google for Developers, OR-Tools documentation. Accessed 17 July 2026 (2024). https://developers.google.com/optimization/cp/cp_solver